

CSC420 Assignment1

Zilun Zhang 1001175685

September 25 2017

1

(1) Load the Boston housing data from the sklearn datasets module

Please see the code file "q1.py". In order to let test result repeatable, I will use the result of setting random seed equals to 0. I will comment the line which sets random seed equals to 0 for the hand in version. If you want to repeat the result, uncomment it will be fine.

(2) Describe and summarize the data in terms of number of data points, dimensions, target, etc

Please see the result of code file "q1.py". I use pandas package to summarize data for each set (X, y) .

```

Sample Of X Summary
      0      1      2      3      4      5  \
count 506.000000 506.000000 506.000000 506.000000 506.000000 506.000000
mean   3.593761 11.363636 11.136779  0.069170  0.554695  6.284634
std    8.596783 23.322453  6.860353  0.253994  0.115878  0.702617
min    0.006320  0.000000  0.460000  0.000000  0.385000  3.561000
25%    0.082045  0.000000  5.190000  0.000000  0.449000  5.885500
50%    0.256510  0.000000  9.690000  0.000000  0.538000  6.208500
75%    3.647423 12.500000 18.100000  0.000000  0.624000  6.623500
max    88.976200 100.000000 27.740000  1.000000  0.871000  8.780000

      6      7      8      9     10     11  \
count 506.000000 506.000000 506.000000 506.000000 506.000000 506.000000
mean  68.574901  3.795043  9.549407 408.237154 18.455534 356.674032
std   28.148861  2.105710  8.707259 168.537116  2.164946  91.294864
min    2.900000  1.129600  1.000000 187.000000 12.600000  0.320000
25%   45.025000  2.100175  4.000000 279.000000 17.400000 375.377500
50%   77.500000  3.207450  5.000000 330.000000 19.050000 391.440000
75%   94.075000  5.188425 24.000000 666.000000 20.200000 396.225000
max  100.000000 12.126500 24.000000 711.000000 22.000000 396.900000

      12
count 506.000000
mean  12.653063
std    7.141062
min    1.730000
25%    6.950000
50%   11.360000
75%   16.955000
max   37.970000
-----
Sample Of y Summary
      0
count 506.000000
mean  22.532806
std    9.197104
min    5.000000
25%   17.025000
50%   21.200000
75%   25.000000
max   50.000000
-----

```

To look at the weight fairly, I normalized data.

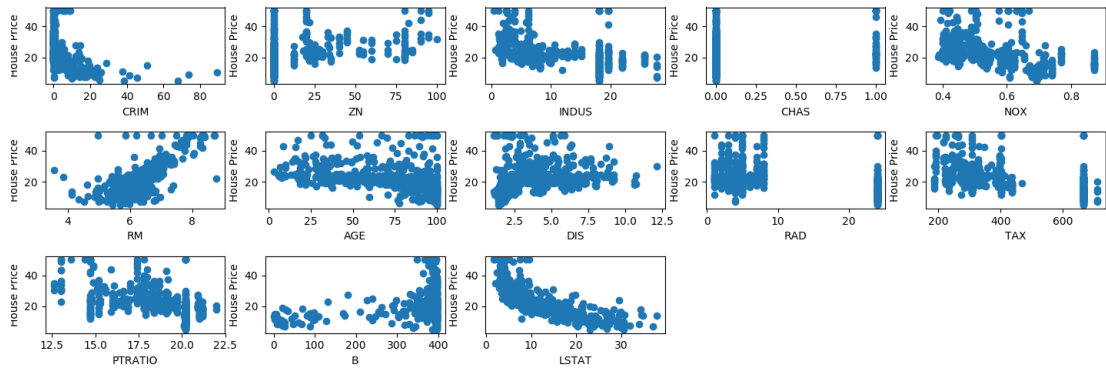
Here is the data summary **after normalization**:

Sample Of X Summary					
	0	1	2	3	4 \
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	6.340997e-17	-6.343191e-16	-2.682911e-15	4.701992e-16	2.490322e-15
std	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00
min	-4.177134e-01	-4.877224e-01	-1.557842e+00	-2.725986e-01	-1.465882e+00
25%	-4.088961e-01	-4.877224e-01	-8.676906e-01	-2.725986e-01	-9.130288e-01
50%	-3.885818e-01	-4.877224e-01	-2.110985e-01	-2.725986e-01	-1.442174e-01
75%	6.248255e-03	4.877224e-02	1.015999e+00	-2.725986e-01	5.986790e-01
max	9.941735e+00	3.804234e+00	2.422565e+00	3.668398e+00	2.732346e+00
	5	6	7	8	9 \
count	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02	5.060000e+02
mean	-1.145230e-14	-1.407855e-15	9.210902e-16	5.441409e-16	-8.868619e-16
std	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00	1.000990e+00
min	-3.880249e+00	-2.335437e+00	-1.267069e+00	-9.828429e-01	-1.313990e+00
25%	-5.686303e-01	-8.374480e-01	-8.056878e-01	-6.379618e-01	-7.675760e-01
50%	-1.084655e-01	3.173816e-01	-2.793234e-01	-5.230014e-01	-4.646726e-01
75%	4.827678e-01	9.067981e-01	6.623709e-01	1.661245e+00	1.530926e+00
max	3.555044e+00	1.117494e+00	3.960518e+00	1.661245e+00	1.798194e+00
	10	11	12		
count	5.060000e+02	5.060000e+02	5.060000e+02		
mean	-9.205636e-15	8.163101e-15	-3.370163e-16		
std	1.000990e+00	1.000990e+00	1.000990e+00		
min	-2.707379e+00	-3.907193e+00	-1.531127e+00		
25%	-4.880391e-01	2.050715e-01	-7.994200e-01		
50%	2.748590e-01	3.811865e-01	-1.812536e-01		
75%	8.065758e-01	4.336510e-01	6.030188e-01		
max	1.638828e+00	4.410519e-01	3.548771e+00		
Sample Of y Summary					
	0				
count	506.000000				
mean	22.532806				
std	9.197104				
min	5.000000				
25%	17.025000				
50%	21.200000				
75%	25.000000				
max	50.000000				

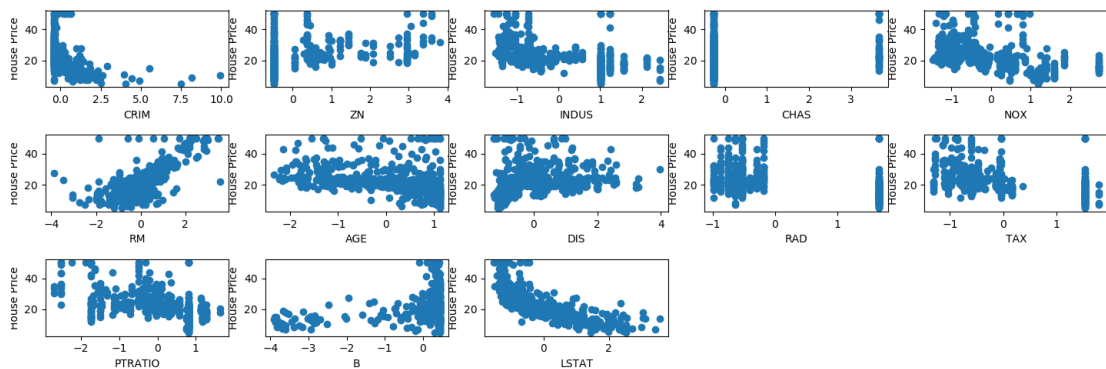
Here is the code of normalization:

```
1 def normalize_data(X):
2     matrix_X_subtract_mean = X - np.mean(X, axis=0, keepdims=True)
3     matrix_X_std = np.power(help_var(X), 0.5)
4     normal_X = matrix_X_subtract_mean/matrix_X_std
5     return normal_X
6
7
8 def cal_var_vector(g):
9     mean = g.mean(0)
10    var_vector = np.zeros((1, g.shape[1]), dtype=float)
11    for i in range(g.shape[0]):
12        var_vector += np.power(g[i, :] - mean, 2)
13    result = np.divide(var_vector, (g.shape[0]))
14    return result
15
16
17 # calculate var along axis 0 for matrix.
18 def help_var(X):
19     # loop features
20     var_list = []
21     for feature in range(X.shape[1]):
22         column = X[:, feature]
23         column = column.reshape(506, 1)
24         var = cal_var_vector(column)
25         var_list.append(var)
26     return np.array(var_list).reshape(1, X.shape[1])
```

(3) **Visualization:** present a single grid containing plots for each feature against the target. Choose the appropriate axis for dependent vs. independent variables.



Here is the scatter graph of data **after** normalization.



Code:

```
1 def visualize(X, y, features):
2     plt.figure(figsize=(20, 5))
3     feature_count = X.shape[1]
4     # i: index
5     for i in range(feature_count):
6         plt.subplot(3, 5, i + 1)
7         # Plot feature i against y
8         plt.scatter(X[:, i], y)
9         plt.xlabel(features[i])
10        plt.ylabel("House Price")
11    plt.tight_layout()
12    plt.show()
```

(4) Divide your data into training and test sets, where the training set consists of 80 % of the data points(chosen at random).

Code:

```
1 def split(X, y, train_split_rate):
2     feature_count, data_X_count, train_X_count = X.shape[1], X.
3     shape[0], ceil(X.shape[0] * train_split_rate)
4     train_X = []
5     test_X = []
6     train_y = []
7     test_y = []
8     training_index = np.random.choice(data_X_count, int(
9     train_X_count), replace=False)
10    for index in range(data_X_count):
11        if index in training_index:
12            train_X.append(X[index])
13            train_y.append(y[index])
14        else:
15            test_X.append(X[index])
16            test_y.append(y[index])
17
18    train_X = np.array(train_X)
19    train_y = np.array(train_y)
20    test_X = np.array(test_X)
21    test_y = np.array(test_y)
22    print("train set of X is: " + str(train_X.shape))
23    print("train set of y is: " + str(train_y.shape))
24    print("test set of X is: " + str(test_X.shape))
25    print("test set of y is: " + str(test_y.shape))
26    print()
27    return train_X, train_y, test_X, test_y
```

(5) Write code to perform linear regression to predict the targets using the training data. Remember to add a bias term to your model.

Code:

```
1 def fit_regression(X,Y):
2     # implement linear regression
3     # Remember to use np.linalg.solve instead of inverting!
4     # add bias term
5     feature_count = X.shape[1]
6     X = np.insert(X, feature_count, 1, axis=1)
7     # (X^T*X)W = X^T*y
8     left = np.dot(X.T, X)
9     right = np.dot(X.T, Y)
10    # solve w
11    return np.linalg.solve(left, right)
```

```
1 def get_predict_value(w, X):
2     feature_count = X.shape[1]
3     # add bias term
4     X = np.insert(X, feature_count, 1, axis=1)
5     return np.dot(X, w)
```

(6) Tabulate each feature along with its associated weight and present them in a table. Explain what the sign of the weight means in the third column ('INDUS') of this table. Does the sign match what you expected? Why?

Before normalization:

CRIM	-0.109703
ZN	0.041546
INDUS	0.010951
CHAS	1.935669
NOX	-17.866761
RM	3.280849
AGE	0.004475
DIS	-1.385640
RAD	0.366823
TAX	-0.015388
PTRATIO	-0.894088
B	0.008872
LSTAT	-0.553587
BIAS	39.570475


The sign of 'INDUS' is positive, which is: 0.010951, which makes sense because more proportion of non-retail business acres per town means less business acres per town, which let business acres become more valued, therefore the retail housing price will increase."

After normalization:


CRIM	-0.942165
ZN	0.968004
INDUS	0.075057
CHAS	0.491162
NOX	-2.068312
RM	2.302902
AGE	0.125840
DIS	-2.914872
RAD	3.190866
TAX	-2.590951
PTRATIO	-1.933739
B	0.809196
LSTAT	-3.949287
BIAS	22.540751

below is a graph of weights against indices:

Before normalization:



After normalization:



(7) Test the fitted model on your test set and calculate the Mean Square Error of the result.

MSE of model by using test data before normalization is: 16.5753061192.

MSE of model by using test data **after normalization** is: 16.5753061192.

Code:

```
1 def mse(predict_value , test_value):
2     # print(predict_value , test_value)
3     return np.mean(np.power(test_value-predict_value , 2))
```

(8) Suggest and calculate more error measurement metrics; justify your choice.

Suggested norm 1 error(abs value), r square value and RMSE error, which are implemented on q1.py file.

Norm 1 loss of model before normalization is: 0.02999964327.

Norm 1 loss of model **after normalization** is: 0.02999964327.

R square coefficient of model before normalization is: 0.83951752255

R square coefficient of model **after normalization** is: 0.83951752255

RMSE loss before normalization is: 4.071278192308554.

RMSE loss **after normalization** is: 4.0712781923084265.

Code:

```
1 def norm1_loss(test_value , predict_value):
2     result_vector = predict_value - test_value
3     return np.mean(np.abs(result_vector))/test_value.shape[0]

1 def r_square_coeff(y, y_predict):
2     ss_total = cal_var_vector(y.reshape(y.shape[0] , 1))*int(y.shape
3     [0])
4     ss_res = mse(y, y_predict)*int(y.shape[0])
5     return float(1 - np.divide(ss_res , ss_total))

1 def rmse_loss(mse):
2     return sqrt(mse)
```

(9) Feature Selection: Based on your results, what are the most significant features that best predict the price? Justify your answer.

Most significant feature to predict the price is LSTAT after normalization, which makes sense because more % lower status of the population in town means less

quality of population in town, less opportunity that businessmen investment in that area, less infrastructure and less security in that area therefore the housing price will decrease in that area. It has negative effect for housing price, that's why it has a negative weight.

Also, RAD is important, it has the largest positive weight above all features. Index of accessibility to radial highways represent how developed and convenient a region. Obviously, more RAD means this region is more developed, therefore the housing price will go high. It has positive effect for housing price, that is why it has a positive weight.

2

(1)

$$\begin{aligned}
 L(\underline{w}) &= \frac{1}{2} \|A^{\frac{1}{2}}(\underline{y} - X\underline{w})\|^2 + \frac{\lambda}{2} \|\underline{w}\|^2 \\
 &= \frac{1}{2} (A^{\frac{1}{2}}(\underline{y} - X\underline{w}))^T (A^{\frac{1}{2}}(\underline{y} - X\underline{w})) + \frac{\lambda}{2} \underline{w}^T \underline{w} \\
 &= \frac{1}{2} (A^{\frac{1}{2}}\underline{y} - A^{\frac{1}{2}}X\underline{w})^T (A^{\frac{1}{2}}\underline{y} - A^{\frac{1}{2}}X\underline{w}) + \frac{\lambda}{2} \underline{w}^T \underline{w} \\
 &= \frac{1}{2} (\underline{y}^T A^{\frac{1}{2}} - \underline{w}^T X^T A^{\frac{1}{2}}) (A^{\frac{1}{2}}\underline{y} - A^{\frac{1}{2}}X\underline{w}) + \frac{\lambda}{2} \underline{w}^T \underline{w} \\
 &= \frac{1}{2} (\underline{y}^T A^{\frac{1}{2}} \cdot A^{\frac{1}{2}}\underline{y} - \underline{y}^T A^{\frac{1}{2}} A^{\frac{1}{2}} X\underline{w} \\
 &\quad - \underline{w}^T X^T A^{\frac{1}{2}} A^{\frac{1}{2}} \underline{y} + \underline{w}^T X^T A^{\frac{1}{2}} A^{\frac{1}{2}} X\underline{w}) \\
 &\quad + \frac{\lambda}{2} \underline{w}^T \underline{w} \\
 &= \frac{1}{2} (\underline{y}^T A \underline{y} - \underline{y}^T A X \underline{w} - \underline{w}^T X^T A \underline{y} + \underline{w}^T X^T A X \underline{w}) \\
 &\quad + \frac{\lambda}{2} \underline{w}^T \underline{w} \\
 &= \frac{1}{2} (\underline{y}^T A \underline{y} + \underline{w}^T X^T A X \underline{w} - 2 \underline{w}^T X^T A \underline{y}) + \frac{\lambda}{2} \underline{w}^T \underline{w} \\
 \frac{\partial L}{\partial \underline{w}} &= \frac{1}{2} (2 X^T A X \underline{w} - 2 X^T A \underline{y}) + \lambda \underline{w} \\
 &= X^T A X \underline{w} - X^T A \underline{y} + \lambda \underline{w}
 \end{aligned}$$

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial \underline{w}} = 0 &\Rightarrow \underline{X}^T \underline{A} \underline{X} \underline{w} + \lambda \underline{w} = \underline{X}^T \underline{A} \underline{y} \\ (\underline{X}^T \underline{A} \underline{X} + \lambda \underline{I}) \underline{w} &= \underline{X}^T \underline{A} \underline{y} \\ \underline{w}^* &= (\underline{X}^T \underline{A} \underline{X} + \lambda \underline{I})^{-1} \underline{X}^T \underline{A} \underline{y} \end{aligned}$$

(2) Please see the file q2.py

Code:

```

1 def LRLS(test_datum, x_train, y_train, tau, lam=1e-5):
2     '''
3     Input: test_datum is a dx1 test vector
4           x_train is the N_train x d design matrix
5           y_train is the N_train x 1 targets vector
6           tau is the local reweighting parameter
7           lam is the regularization parameter
8     output is y_hat the prediction on test_datum
9     '''
10
11     # (X^T * A * X + lamda * I) W = X^T * A * Y
12     # build a empty list contain diagonal of A.
13     x_train_N = x_train.shape[0]
14     x_train_d = x_train.shape[1]
15     A = np.zeros((x_train_N, x_train_N), dtype=float)
16     identity_matrix = np.identity(x_train_d)
17     content_numerator = -l2(x_train, test_datum.T)
18     content_denominator = np.multiply(2, np.power(tau, 2))
19     numerator = np.divide(content_numerator, content_denominator)
20     log_summation = misc.logsumexp(numerator)
21     denominator = np.exp(log_summation)
22
23     for i in range(x_train_N):
24         A[i, i] = np.exp(numerator[i]) / denominator
25
26     temp = np.matmul(x_train.T, A)
27     left_temp = np.matmul(temp, x_train)
28     left = left_temp + np.multiply(lam, identity_matrix)
29     right = np.dot(temp, y_train)
30     w = np.linalg.solve(left, right)
31     y_hat = np.matmul(test_datum.T, w)
32     return float(y_hat)

```

```

1 def run_k-fold(x, y, taus, k):
2     '''


```

```

3 Input: x is the N x d design matrix
4       y is the N x 1 targets vector
5       taus is a vector of tau values to evaluate
6       K in the number of folds
7 output is losses a vector of k-fold cross validation losses one
8       for each tau value
9
10 loss = []
11 # random process
12 concatenate_matrix = np.concatenate((x, y[:, None]), axis=1)
13 np.random.shuffle(concatenate_matrix)
14 print(concatenate_matrix[0, 0])
15 # split to k fold
16 temp_container = np.array_split(concatenate_matrix, k)
17 i = 0
18 while i < k:
19     print("fold number is:")
20     print(i+1)
21     fold = np.array(temp_container[i])
22     x_test = fold[:, :d]
23     print("x_test size is:")
24     print(x_test.shape)
25     y_test = fold[:, d]
26     print("y_test size is:")
27     print(y_test.shape)
28     x_train = []
29     y_train = []
30     for j in range(k):
31         if j == i:
32             pass
33         else:
34             other_fold = np.array(temp_container[j])
35             print("other_fold shape is:")
36             print(other_fold.shape)
37             x_train.append(other_fold[:, :d])
38             y_train.append(other_fold[:, d])
39     i += 1
40
41 x_train = np.concatenate(np.array(x_train))
42 y_train = np.concatenate(np.array(y_train))
43
44 print("x train shape:")
45 print(x_train.shape)
46 print("y train shape:")
47 print(y_train.shape)
48 print("x test shape:")
49 print(x_test.shape)
50 print("y test shape:")
51 print(y_test.shape)
52 temp_loss = run_on_fold(x_test, y_test, x_train, y_train,
53 tau)
54 loss.append(temp_loss)
55 output = np.array(loss).mean(0)
56 print(output.shape)
57 print(output)
58 return output

```


(3)



(4)

When $r \rightarrow \infty$, the loss converges around 22.4. When $r \rightarrow 0$, the loss goes to infinity.

For $r = [10, 10^{10}]$:



3

(1)

$$S = \{a_1, \dots, a_n\}.$$

$$\text{LHS} = E_I \left[\frac{1}{m} \sum_{i \in I} a_i \right]$$

$$= \frac{1}{m} \sum_{i=1}^m E(a_i)$$

$$= \frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n \frac{1}{n} a_j$$

$$= \frac{1}{n} \sum_{j=1}^n a_j$$

$$= \frac{1}{n} \sum_{i=1}^n a_i$$

$$= \text{RHS}.$$

$$E(x) = \sum_{i=1}^n x_i P(x_i)$$

since each index
is drawn uniformly
from the set

$$\{1, \dots, n\} \rightarrow \frac{1}{n}$$

(2)

$$\text{know: } L_I(\mathbf{x}, \mathbf{y}, \theta) = \frac{1}{m} \sum_{i=1}^m L(\mathbf{x}_i^{(i)}, \mathbf{y}_i^{(i)}, \theta)$$

$$\Rightarrow \nabla L_I(\mathbf{x}, \mathbf{y}, \theta) = \frac{1}{m} \sum_{i=1}^m \nabla L(\mathbf{x}_i^{(i)}, \mathbf{y}_i^{(i)}, \theta)$$

(differentiate both side preserved linearity.)

$$E_I [\nabla L_I(\mathbf{x}, \mathbf{y}, \theta)] = E_I \left[\frac{1}{m} \sum_{i=1}^m \nabla L_i(\mathbf{x}_i, \mathbf{y}_i, \theta) \right]$$

$$\begin{aligned} E(\mathbf{x}) &= \sum_{i=1}^n \mathbf{x}_i P(\mathbf{x}_i) \\ &= \frac{1}{m} \sum_{i=1}^m E(\nabla L_i(\mathbf{x}_i, \mathbf{y}_i, \theta)) \\ &= \frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n \frac{1}{n} \nabla L_j(\mathbf{x}_j, \mathbf{y}_j, \theta) \\ &= \frac{1}{n} \sum_{j=1}^n \nabla L_j(\mathbf{x}_j, \mathbf{y}_j, \theta) \\ &= \nabla L(\mathbf{x}, \mathbf{y}, \theta) \\ &= \text{RHS.} \end{aligned}$$

since each index is drawn uniformly from the set $\{1, \dots, n\} \rightarrow \frac{1}{n}$

(3) It shows that in SGD, expected value of a mini-batch's loss, $E(\nabla L_I(\mathbf{x}, \mathbf{y}, \theta))$ is equal to the true empirical gradient of loss over whole data set $\nabla L(\mathbf{x}, \mathbf{y}, \theta)$ (a mini-batch's loss is an unbiased estimator to evaluate true gradient of loss of whole data set).

(4)

(a)

$$\frac{\partial L}{\partial \omega} = \frac{-2\mathbf{x}^T(\mathbf{y} - \mathbf{x}\omega)}{M} \quad (1)$$

(b)

```
1 def lin_reg_gradient(X, y, w):
2     """
3     Compute gradient of linear regression model parameterized by w
4     """
5     s = X.shape[0]
6     gradient_w = - 2 * np.matmul(X.T, (y - np.matmul(X, w)))/s
7     return gradient_w
```

(5) In order to let the result repeatable, I set `np.random.seed` to be 0 again. Comment it if you need random value.

true gradient is:

```
true gradient is:  
[ 7.48161897e+02  2.66119727e+03  9.20005630e+03  8.36869099e+03  
 5.19727844e+01  4.13891293e+02  4.69502078e+03  5.09428703e+04  
 2.87367326e+03  7.18470104e+03  3.07658588e+05  1.38303337e+04  
 2.76270255e+05  9.42091869e+03]
```

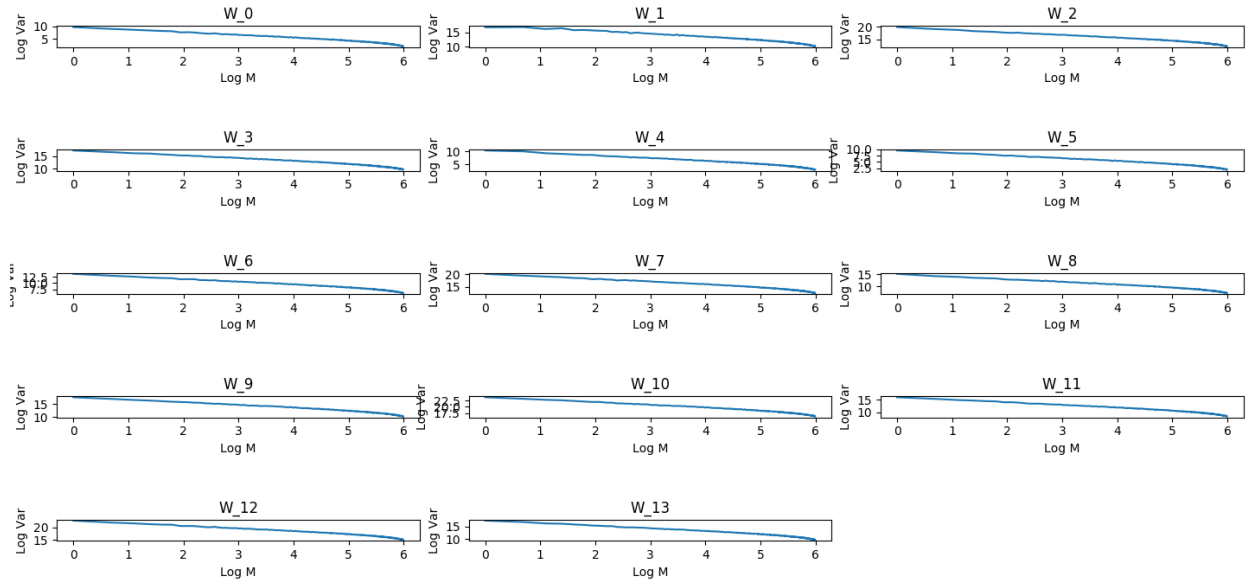
MSE is: 43870.8653703

Cosine similarity loss is: 0.999998706366

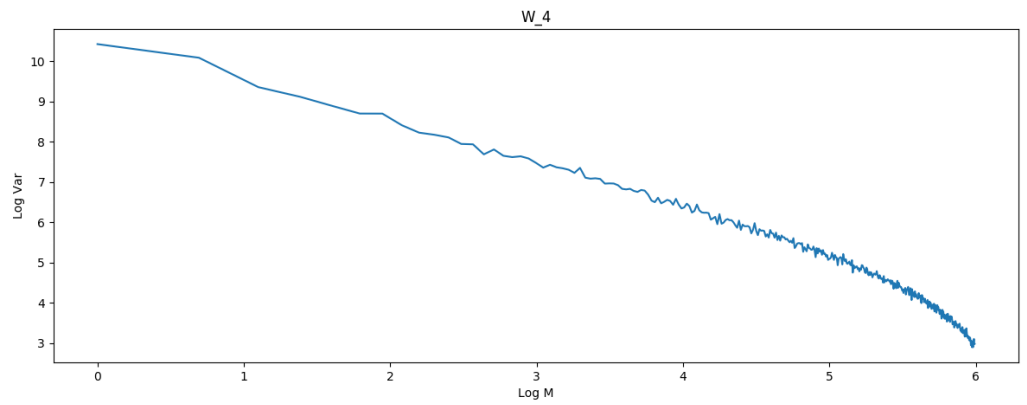
I think they are both meaningful, but in this case, cosine similarity is more meaningful. The reason is initial magnitude for each element in gradient vector are too large. So a small variation between prediction value and test value for each element in gradient vector will cause a huge difference, not mention square and summation all of them when we calculate MSE. Since vectors we compare has 14 dimensions, look at the angle between them is a good idea to measure the error in above situation (each entry's magnitude is too large). Also, the step we are doing is the first step of SGD, since the batch is randomly chosen, it can be far away from the optimal position in the beginning, therefore, for the ultimate purpose, a right direction is more important than MSE.

(6)

Result:



For a single w , i.e. w_4 (start from w_0):



Code:

```
1 def calculate_gradient_var(batch_sampler, w, K):
2     i = 0
3     # init a 500 * 14 matrix
4     gradient_matrix = np.zeros((K, batch_sampler.features), dtype =
5         float)
6     total_loss_gradient = 0
7     while i < K:
8         X_b, y_b = batch_sampler.get_batch()
9         batch_grad = lin_reg_gradient(X_b, y_b, w)
10        gradient_matrix[i, :] = batch_grad
11        total_loss_gradient += batch_grad
12        i += 1
13    gradient = np.divide(total_loss_gradient, K)
14    # take variance for each column result is a (1,14) matrix
15    # vector_variance = np.log(np.sqrt(gradient_matrix.var(0)))
16    vector_variance = np.log(cal_var(gradient_matrix))
17    print("shape of vector variance is:")
18    print(vector_variance.shape)
19    return gradient, vector_variance
20
21 def cal_var(g):
22     mean = g.mean(0)
23     var_vector = np.zeros((1, g.shape[1]), dtype=float)
24     for i in range(g.shape[0]):
25         var_vector += np.power(g[i, :] - mean, 2)
26     result = np.divide(var_vector, (g.shape[0]))
27     return result
28
29
30 def plot_delta_j_vs_m(X, y, w, m_start, m_end, K):
31     m_range = m_end - m_start
32     x_list = np.arange(1, m_range + 1)
33     x_list = np.log(x_list)
34     x_list.tolist()
35     variance_matrix = np.zeros((m_range, X.shape[1]), dtype=float)
36     plt.figure(figsize=(20, 5))
37     for i in range(1, m_range + 1):
38         batch_sampler = BatchSampler(X, y, i)
39         gradient, vector_variance = calculate_gradient_var(
40             batch_sampler, w, K)
41         print("i is:")
42         print(i)
43         variance_matrix[i - 1, :] = vector_variance
44
45     for j in range(X.shape[1]):
46         print("j is: ")
47         print(j)
48         # y to plot is jth column for the vector variance.
49         print(variance_matrix.shape)
50         y_list = variance_matrix[:, j]
51         plt.subplot(4, 4, j + 1)
52         plt.plot(x_list, y_list)
```

```
52     plt.title("W_{}".format(j))
53     plt.xlabel("Log M")
54     plt.ylabel("Log Var")
55     plt.tight_layout()
56     plt.show()
```